

① CREATE TABLE Athletes
{
Name VARCHAR(255)
DOB CHAR(10),
PRIMARY KEY (Name, DOB)
};

CREATE TABLE Sports
{
Name VARCHAR(255), primary key,
ISolympic BIT,
FOREIGN KEY (Name) references (Name)
};

In these create table statements I made the foreign key 'name' reference the pk 'name' from the athlete table.

② $R(ABCD)$
 $F = \{A \rightarrow B, B \rightarrow D, A \rightarrow D\}$

(a) True. Because $B \rightarrow D$, we can also say that $D \rightarrow B$. Therefore $A \rightarrow DC$, which is the same as $A \rightarrow CD$, is valid.

(b) True.

$A^+ = ABCD$

(c) $F_c = \{$
 $A \rightarrow B$
 $B \rightarrow D$
 $\}$

Yes, this is redundant because $A \rightarrow D$ is already implied due to $A \rightarrow B$ and $B \rightarrow D$.

③ (a) There are 3 candidate keys: AB, AC, AD

(b) Yes, R is in 3NF because FDs $AB \rightarrow C$; $AC \rightarrow D$ have a ck on the left and the last FD has part of {AB} on the right. FD $D \rightarrow B$.

(c) ABCD is decomposed into $R_1(\underline{ACD})$, with the keys AC, AD, and the FDs $AC \rightarrow D$, $AD \rightarrow C$, and $R_2(\underline{DB})$ with key D and FD $D \rightarrow B$. Both $R_1 + R_2$ are in BCNF so decomposition ends.

④ (a) I would use a hash index on id. The index is on the pk so at most one record can match. If a match is found then it will take 1 page from disk.

(b) I would use the unclustered B+-tree index on (age, rating). The leaves will have the data that we query for and it won't take pages from disk.